

Task 4 Machine Learning Internship: Haris Khan

```
[63]: import tensorflow as tf
      from tensorflow.keras import datasets, layers, models
      import matplotlib.pyplot as plt
      import numpy as np

[64]: (X_train, y_train), (X_test, y_test) = datasets.cifar10.load_data()
      X_train.shape

[64]: (50000, 32, 32, 3)

[65]: X_test.shape

[65]: (10000, 32, 32, 3)

[66]: y_train.shape

[66]: (50000, 1)

[67]: y_train[:5]

[67]: array([[6],
          [9],
          [9],
          [4],
          [1]], dtype=uint8)

[68]: y_train = y_train.reshape(-1,)
      y_train[:5]


[68]: array([6, 9, 9, 4, 1], dtype=uint8)

[69]: y_test = y_test.reshape(-1,)


[70]: classes = ["airplane", "automobile", "bird", "cat", "deer", "dog", "frog", "horse", "ship", "truck"]

[71]: def plot_sample(X, y, index):
      plt.figure(figsize = (15,2))
      plt.imshow(X[index])
      plt.xlabel(classes[y[index]])

[72]: plot_sample(X_train, y_train, 0)
```



```
[73]: plot_sample(X_train, y_train, 1)
```



```
[74]: X_train = X_train / 255.0
      X_test = X_test / 255.0
```

```
[75]: ann = models.Sequential([
      layers.Flatten(input_shape=(32,32,3)),
      layers.Dense(3000, activation='relu'),
      layers.Dense(1000, activation='relu'),
      layers.Dense(10, activation='softmax')
      ])

      ann.compile(optimizer='SGD',
                  loss='sparse_categorical_crossentropy',
                  metrics=['accuracy'])

      ann.fit(X_train, y_train, epochs=5)

/home/haris-khan/jupyter-env/lib/python3.12/site-packages/keras/src/layers/reshaping/flatten.py:37: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.
  super().__init__(**kwargs)
Epoch 1/5
1563/1563 — 60s 38ms/step - accuracy: 0.3028 - loss: 1.9363
Epoch 2/5
1563/1563 — 59s 38ms/step - accuracy: 0.4193 - loss: 1.6480
Epoch 3/5
1563/1563 — 60s 39ms/step - accuracy: 0.4546 - loss: 1.5531
Epoch 4/5
1563/1563 — 63s 40ms/step - accuracy: 0.4702 - loss: 1.4992
Epoch 5/5
1563/1563 — 62s 40ms/step - accuracy: 0.4926 - loss: 1.4468
```

```
[75]: <keras.src.callbacks.history.History at 0x7c0ca3755460>

[41]: from sklearn.metrics import confusion_matrix, classification_report
      import numpy as np
      y_pred = ann.predict(X_test)
      y_pred_classes = [np.argmax(element) for element in y_pred]
```

```
[41]: from sklearn.metrics import confusion_matrix, classification_report
import numpy as np
y_pred = ann.predict(X_test)
y_pred_classes = [np.argmax(element) for element in y_pred]

print("Classification Report: \n", classification_report(y_test, y_pred_classes))
```

```
313/313 ————— 3s 10ms/step
Classification Report:
      precision    recall  f1-score   support

     0       0.51       0.57       0.54       1000
     1       0.74       0.41       0.53       1000
     2       0.47       0.16       0.24       1000
     3       0.32       0.31       0.32       1000
     4       0.63       0.12       0.21       1000
     5       0.32       0.48       0.38       1000
     6       0.41       0.71       0.52       1000
     7       0.48       0.63       0.54       1000
     8       0.55       0.69       0.62       1000
     9       0.56       0.58       0.57       1000

 accuracy          0.47       10000
 macro avg          0.50       0.45       10000
 weighted avg       0.50       0.45       10000
```

```
[42]: cnn = models.Sequential([
    layers.Conv2D(filters=32, kernel_size=(3, 3), activation='relu', input_shape=(32, 32, 3)),
    layers.MaxPooling2D((2, 2)),

    layers.Conv2D(filters=64, kernel_size=(3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),

    layers.Flatten(),
    layers.Dense(64, activation='relu'),
    layers.Dense(10, activation='softmax')
])
```

/home/haris-khan/jupyter-env/lib/python3.12/site-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an 'input_shape'/'input_dim' argument to a layer. When using Sequential models, prefer using an 'Input(shape)' object as the first layer in the model instead.
super().__init__(activity_regularizer=activity_regularizer, **kwargs)

```
[43]: cnn.compile(optimizer='adam',
                loss='sparse_categorical_crossentropy',
                metrics=['accuracy'])
```

```
[44]: cnn.fit(X_train, y_train, epochs=10)

Epoch 1/10
1563/1563 ————— 26s 16ms/step - accuracy: 0.3968 - loss: 1.6608
Epoch 2/10
1563/1563 ————— 26s 17ms/step - accuracy: 0.6104 - loss: 1.1177
Epoch 3/10
```

```
[44]: cnn.fit(X_train, y_train, epochs=10)

Epoch 1/10
1563/1563 ————— 26s 16ms/step - accuracy: 0.3968 - loss: 1.6608
Epoch 2/10
1563/1563 ————— 26s 17ms/step - accuracy: 0.6104 - loss: 1.1177
Epoch 3/10
1563/1563 ————— 26s 16ms/step - accuracy: 0.6637 - loss: 0.9710
Epoch 4/10
1563/1563 ————— 26s 16ms/step - accuracy: 0.6951 - loss: 0.8733
Epoch 5/10
1563/1563 ————— 26s 16ms/step - accuracy: 0.7276 - loss: 0.7950
Epoch 6/10
1563/1563 ————— 26s 16ms/step - accuracy: 0.7400 - loss: 0.7422
Epoch 7/10
1563/1563 ————— 26s 17ms/step - accuracy: 0.7629 - loss: 0.6793
Epoch 8/10
1563/1563 ————— 26s 17ms/step - accuracy: 0.7811 - loss: 0.6326
Epoch 9/10
1563/1563 ————— 26s 17ms/step - accuracy: 0.7960 - loss: 0.5829
Epoch 10/10
1563/1563 ————— 26s 17ms/step - accuracy: 0.8096 - loss: 0.5477
```

```
[44]: <keras.src.callbacks.history.History at 0x7c0d09f625d0>
```

```
[45]: y_pred = cnn.predict(X_test)
y_pred[:5]
```

```
313/313 ————— 2s 5ms/step
```

```
[45]: array([[2.5352151e-06, 2.6942054e-05, 1.7831140e-04, 9.6407247e-01,
1.3295193e-04, 2.1179426e-02, 1.1658215e-02, 9.5442022e-07,
2.7371121e-03, 1.1200614e-05],
[3.0432458e-03, 2.3638444e-02, 1.5797217e-04, 8.7096168e-06,
1.4406257e-07, 6.5121561e-08, 3.8977470e-07, 1.1124809e-07,
9.6196413e-01, 1.1186650e-02],
[2.2752885e-02, 9.5861517e-02, 1.6352712e-04, 1.3839080e-03,
4.7590127e-04, 2.4885873e-04, 6.0887010e-05, 1.1321453e-03,
7.9294497e-01, 8.4975407e-02],
[7.8561568e-01, 1.1721703e-02, 1.6631883e-01, 1.0808820e-03,
9.8324716e-03, 1.2070191e-04, 7.9765778e-05, 2.2857585e-04,
2.4680052e-02, 3.2115870e-04],
[1.6112288e-07, 2.5472340e-05, 9.5093315e-03, 1.5235729e-02,
1.6124251e-01, 1.5780315e-03, 8.1240332e-01, 1.0699318e-06,
4.0293176e-06, 4.0363693e-07]], dtype=float32)
```

```
[46]: y_classes = [np.argmax(element) for element in y_pred]
y_classes[:5]
```

```
[46]: [3, 8, 8, 0, 6]
```

```
[47]: y_test[:5]
```

```
[47]: array([3, 8, 8, 0, 6], dtype=uint8)
```

```
[48]: nInt sample(X_test, y_test, 3)
```

```
[45]: array([[2.552151e-08, 2.0942054e-05, 1.7851140e-04, 9.040724/e-01,  
1.3295193e-04, 2.1179426e-02, 1.1658215e-02, 9.5442022e-07,  
2.7371121e-03, 1.1200614e-05],  
[3.0432458e-03, 2.3638444e-02, 1.5797217e-04, 8.7096168e-06,  
1.4406257e-07, 6.5121561e-08, 3.8977470e-07, 1.1124809e-07,  
9.6196413e-01, 1.1186650e-02],  
[2.2752885e-02, 9.5861517e-02, 1.6352712e-04, 1.3839080e-03,  
4.7590127e-04, 2.4885873e-04, 6.0887010e-05, 1.1321453e-03,  
7.9294497e-01, 8.4975407e-02],  
[7.8561568e-01, 1.1721703e-02, 1.6631883e-01, 1.0808820e-03,  
9.8324716e-03, 1.2070191e-04, 7.9765778e-05, 2.2857585e-04,  
2.4680052e-02, 3.2115870e-04],  
[1.6112288e-07, 2.5472340e-05, 9.5093315e-03, 1.5235729e-02,  
1.6124251e-01, 1.5780315e-03, 8.1240332e-01, 1.0699318e-06,  
4.0293176e-06, 4.0363693e-07]], dtype=float32)
```


```
[46]: y_classes = [np.argmax(element) for element in y_pred]  
y_classes[:5]
```

```
[46]: [3, 8, 8, 0, 6]
```

```
[47]: y_test[:5]
```

```
[47]: array([3, 8, 8, 0, 6], dtype=uint8)
```

```
[48]: plot_sample(X_test, y_test,3)
```



```
[49]: classes[y_classes[3]]
```

```
[49]: 'airplane'
```

```
[50]: classes[y_classes[3]]
```

```
[50]: 'airplane'
```

```
[ ]:
```